# An Unknown Signal Report

George Herbert cj19328@bristol.ac.uk

April 19, 2021

#### Abstract

This report demonstrates my understanding of the methods I have used, the results I have obtained, and my understanding of issues such as overfitting for the 'An Unknown Signal' coursework.

### 1 Equations for linear regression

For a set of (x, y) coordinates that lie along a line with Gaussian noise, with the relationship  $\mathbf{y} = \mathbf{X}\mathbf{w} + \epsilon$  where  $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$ , the maximum likelihood estimation of  $\mathbf{w}$  is equivalent to the least square error estimation and is given by the equation:

$$\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}.$$

This equation is implemented in my code as the following method:

def regressionNormalEquation(self, X, y):
return np.linalg.inv(X.T @ X) @ X.T @ y

X can take one of the following three forms:

$$\mathbf{X} = \begin{bmatrix} x_1 & 1 \\ \vdots & \vdots \\ x_{20} & 1 \end{bmatrix} \quad \mathbf{X} = \begin{bmatrix} x_1^n & x_1^{n-1} & \dots & 1 \\ \vdots & \vdots & \ddots & \vdots \\ x_{20}^n & x_{20}^{n-1} & \dots & 1 \end{bmatrix} \quad \mathbf{X} = \begin{bmatrix} f(x_1) & 1 \\ \vdots & \vdots \\ f(x_{20}) & 1 \end{bmatrix}$$

depending on whether the line is linear, polynomial of degree n, or the unknown function f, respectively.

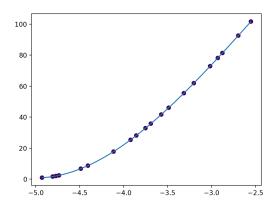
### 2 Choice of polynomial degree

Having created a program, 'display.py', to visualise the points, I drew up a list of segments that appeared nonlinear. Then, I created a program, 'degree.py', that calculated the cross-validation error for these line segments when trained using a model with a polynomial of degree 2 to a polynomial of degree 10; Table 1 shows a small section from the output of this program.

Having analysed the output, it was clear that a large proportion of the nonlinear signals had their minimum cross-validation error when fitted with a polynomial of degree 3. This

Table 1: Section of the output from 'degree.py'

Filename	Line segment	Polynomial degree	Cross-validation error
basic_3.csv	0	2	7.3947610358752875
$basic_3.csv$	0	3	$1.2989585613760917 \mathrm{e}\text{-}23$
:	:	:	:
$adv\_3.csv$	5	9	318.8443359827487
$adv_3.csv$	5	10	279.2750683133305



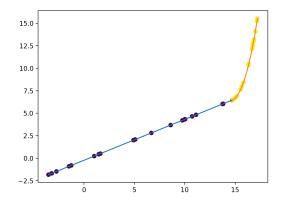


Figure 1: Plot for 'basic\_3.csv'

Figure 2: Plot for 'basic\_4.csv'

consistent minimum cross-validation error indicated that the polynomial line segments in the unknown signal are cubic.

Having incorporated cubic regression into my 'lsr.py' least-squares regression program, I ran the program on the 'basic\_3.csv' and 'basic\_4.csv' unknown signals; Figure 1 and Figure 2 display these plots, respectively. The lines being a near-perfect fit allowed me to validate that a cubic polynomial is reasonable visually.

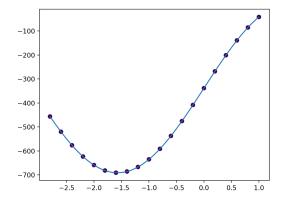
#### 3 Choice of unknown function

Using my 'display.py' program to visualise the signals, I produced a list of potential 'unknown functions' that could represent the underlying signal of line segments, based on their shapes:  $\mathbf{w}_1 sin(x) + \mathbf{w}_2$ ,  $\mathbf{w}_1 cos(x) + \mathbf{w}_2$ ,  $\mathbf{w}_1 tan(x) + \mathbf{w}_2$  and  $\mathbf{w}_1 e^x + \mathbf{w}_2$ .

I then created a program, 'unknown.py', that calculated the cross-validation error for each of the nonlinear line segments previously identified when trained using each of the potential unknown functions. A table displayed the cross-validation errors—similar to that used to determine the polynomial degree.

Having analysed the cross-validation errors, it was clear that all nonlinear signals that were likely not a cubic polynomial had their minimum cross-validation error when trained to fit the function  $\mathbf{w}_1 sin(x) + \mathbf{w}_2$ .

Having incorporated regression to fit a function of the form  $\mathbf{w}_1 sin(x) + \mathbf{w}_2$  into my 'lsr.py' least-squares regression program, I ran the program on the 'basic\_5.csv' and 'adv\_3.csv' unknown signals; the outputs are shown in Figure 3 and Figure 4 respectively. The lines being a near-perfect fit allowed me to validate that the unknown function is of



3000 - 2500 - 2000 - 1500 - 1000 - 500 - - - 500 - - - 1000 - - 1000 - - 1000 - - 1000

Figure 3: Plot for 'basic\_5.csv'

Figure 4: Plot for 'adv\_3.csv'

the form  $\mathbf{w}_1 sin(x) + \mathbf{w}_2$  visually.

#### 4 Model selection

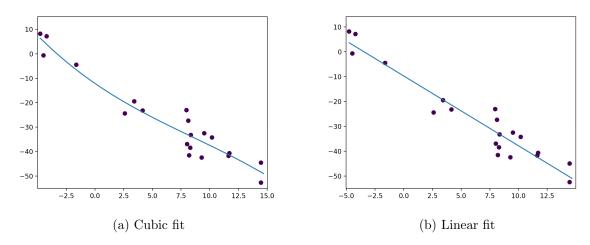


Figure 5: Two fits for the first line segment in 'noise\_2.csv'

Overfitting occurs when a machine learning algorithm produces a model that has learnt the noise in the data as if it represents the structure of the underlying model [1]. In the case of linear regression, overfitting is most likely to occur by producing a model with too complex a function type, such that it would fail to predict future observations. Figure 5 shows an example of this. On the left, (a) is a cubic fit that is too complex a function type for the data points. It would not predict future y-values for a set of given x-values as successfully as the linear fit (b) on the right.

To prevent overfitting, I have used leave-one-out cross-validation when producing a model for each 20-point line segment. Leave-one-out cross-validation is an extreme case of k-fold cross validation such that k=n, where n is the number of data points (in this case, 20). Despite being computationally expensive, I believe that leave-one-out cross-validation is an appropriate technique to prevent overfitting in this case, owing to the limited sample size of each line segment.

Leave-one-out cross-validation involves using each of the 20 data points exactly once as validation data for a model trained using the other 19 data points. The cross-validation error for each function type is calculated as follows [2]:

$$CV_{(n)} = \frac{1}{n} \sum_{i=1}^{n} (y_i - \hat{y}^{(-i)})^2$$

where n is the number of data points in a line segment (i.e. 20),  $y_i$  is the actual y-value for the i-th data point, and  $\hat{y}^{(-i)}$  is the predicted y-value for the i-th data point when trained without using the i-th sample.

The function type with the lowest cross-validation error is then selected for each line segment, and the weights  $\hat{\mathbf{w}}$  are determined by training on all data points for that segment.

## 5 Optimisations and improvements

To begin with, computing the matrix inverse using the np.linalg.inv method is computationally expensive and unnecessary. Instead, given X and y, the maximum likelihood estimation of w could be directly computed as follows: np.linalg.solve(X.T @ X, X.T @ y). Computing  $\hat{\mathbf{w}}$  directly would be faster, as np.linalg.inv computes the inverse of a matrix  $\mathbf{A}$  by solving for  $\mathbf{A}^{-1}$  in  $\mathbf{A}\mathbf{A}^{-1} = \mathbf{I}$  [3]. Thus, there would be a performance benefit by solving for  $\hat{\mathbf{w}}$  in  $\mathbf{X}^T\mathbf{X}\hat{\mathbf{w}} = \mathbf{X}^T\mathbf{y}$  directly.

Another computationally expensive operation in my algorithm is that used to calculate the cross-validation error using leave-one-out cross-validation. The method currently involves fitting the model and calculating the sum squared error n times. Instead, there exists a faster method I could have adopted that involves calculating the leverage. Despite this, I opted not to include this method because my program, as it currently stands, can be easily adapted to use k-fold cross-validation for any value of k that is a factor of 20—changing the constant 'K' in the code achieves this.

### 6 Testing

I created a file, 'test.py', that uses the unittest framework to test each of the methods in 'lsr.py'.

#### References

- [1] Burnham, K. P. and Anderson, D. R. (2002) Model Selection and Multimodel Inference. 2nd ed. Springer-Verlag.
- [2] Taylor, J. (2020) Leave one out cross-validation (LOOCV) STATS 202 https://web.stanford.edu/class/stats202/notes/Resampling/LOOCV.html
- [3] Muldal, A. (2017) Why does numpy.linalg.solve() offer more precise matrix inversions than numpy.linalg.inv()? https://stackoverflow.com/a/31257909/8540479