

An Unknown Signal Report

George Herbert
cj19328@bristol.ac.uk

April 1, 2021

Abstract

This report demonstrates my understanding of the methods I have used, the results I have obtained and my understanding of issues such as overfitting for the ‘An Unknown Signal’ coursework.

1 Equations for linear regression

For a set of points that lie along a line with Gaussian noise $\mathbf{y} = \mathbf{X}\mathbf{w} + \epsilon$ where $\epsilon_i \sim \mathcal{N}(0, \sigma^2)$, the maximum likelihood estimation is equivalent to the least square error estimation and is given by the equation:

$$\hat{\mathbf{w}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}.$$

I’ve implemented this equation in my code as the following function:

```
def regressionNormalEquation(self, X, y):  
    return np.linalg.inv(X.T @ X) @ X.T @ y
```

2 Choice of polynomial order

3 Choice of unknown function

4 Overfitting

Overfitting occurs when a machine learning algorithm produces a model that has learnt the noise in the data as if it represents the structure of the underlying model. [1]

In the case of linear regression, overfitting is most likely to occur by producing a model with too complex a function class, such that it would fail to predict future observations.

5 Model selection

To prevent overfitting, I have used leave-one-out cross-validation when producing a model for each 20-point line segment. Leave-one-out cross-validation is equivalent to k -fold cross validation, such that $k = n$, where n is the number of data-points (in this case 20).

To prevent overfitting each line segment, I have used a type of cross-validation known as k -fold cross-validation—in particular, I have used 5-fold cross-validation. I opted to use k -fold cross-validation due to the small number of datapoints in each segment.

To begin with, k -fold cross-validation involves shuffling and splitting each 20-point segment into k evenly-sized subsamples, as implemented in the `getKFold()` method. Then, each subsample is used exactly once as validation data, whilst the other $k - 1$ partitions are used as training data. The cross-validation error for each model is calculated as follows [2]:

$$CV(\hat{f}) = \frac{1}{N} \sum_{i=1}^N SSE(y_i, \hat{f}_i(x_i))$$

where \hat{f} is a given model; \hat{f}_i is the fitted function, trained with the i -th subsample removed; N is the number of folds (in this case 5); SSE is the sum squared error function; y_i is the y datapoints of subsample i ; and x_i is the x datapoints of subsample i .

The model with the lowest cross-validation error is then selected.

6 Testing

7 Optimisations and improvements

1. Using `np.linalg.solve`
2. Not shuffling the data before leave-one-out cross-validation
3. Not finding the average cross-validation error

References

- [1] Burnham, K. P. and Anderson, D. R. (2002) *Model Selection and Multimodel Inference*. 2nd ed. Springer-Verlag.

- [2] Hastie, T., Tibshirani, R. and Friedman, J. (2001) *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. 2nd ed. New York: Springer.